Erasure Coding and Replication for Edge Devices

Sheshadri K R

Department of Computational and Data Sciences
Indian Institute of Science, Bangalore 560012 INDIA
Email: sheshadrik@iisc.ac.in

1. Problem Definition

The study investigates the problem of providing a reliable distributed data store for an edge device deployment that makes use of the collective storage space of the connected edge devices.

In doing so following questions are addressed :

- What is the impact of using erasure coding to store data reliably at the edge devices? Considering that the edge devices have less compute resources, how does storage and compute trade off?
- How does the erasure coding affect the read and write performance? How slow/fast does it get compared to full replication?
- What is the impact of failure of Edge devices and the costs of recovery(time) in such cases? How does it fare against recovery using replication?
- Examine Reed Solomon Coding techniques with varying values of N and K, investigate its impact on storage efficiency and read/write performance. *

The above mentioned questions are validated through building a data store to provide distributed file storage in the Edge devices, whcih supports both replication and erasure coding.

2. Novelty and Scalability

The proposed scheme is called as Partial Replication, Partial erasure coding (PRPE), in this technique, the client which wants to store a file in the data store, will specify an associated **storage space budget**.

Given this space budget and the file, our job is to find out what is the optimal way of utilizing the storage budget by using both erasure coding and replication. The given file is split into a number of smaller blocks of identical size. Depending on the available storage budget, a fraction of the total blocks will be erasure coded and the rest will be replicated.

The good thing about this technique is that, irrespective of what percentage of file is erasure coded or replicated, the reliability of the file remains the same as a whole. The file will be fault tolerant for same number of edge failures despite using a combination of both the techniques.

This study is carried out on Fog and edge devices, where Fog device acts as a master and edge devices are the storage layer, and are managed by the Fog device. Edge devices are moderately resourced compared to the Fog devices.

To put the storage, main memory and compute in perspective for edge devices, we can consider Raspberry Pi 3 as an edge device which comes with a modest 1.2Ghz single core processor, with around 1GB of main memory and 32GBs of SD card storage. Compare this to a desktop computer, which has around 3.4Ghz 8-core processor with 8GB of main memory and 500GB of disk storage (we can call this as Fog).

As far as scalability is concerned, as and when the requests for storing more files come in, adding more edge devices will increase the collective storage space , hence, allowing us to store the files in the data store. The system naturally is expected to scale with more edge devices joining in, the only bottleneck being the amount of main-memory in the master node (Fog device).

Here, it is shown that the PRPE technique works well with 10 devices. Though these are in no way close to the massive real world deployments. It provides a study of the *"compute vs storage"* trade off, which should serve as good indicator which aids in taking educated decisions while building a data store.

3. Goals and Non Goals

Here I outline the minimum goals I would want to achieve.

- Design and develop a data store a central master node and a set of edge devices which are used as storage nodes.
- The reliability of the file as a whole should remain the same despite the techniques used underneath.
- Evaluate the time taken for doing reads and writes for files, by varying the allotted budget of storage.
- Examine the effect of varying the parameters of the Erasure coding (N and K for the Reed Solomon Coding technique) *
- Compare this technique with full replication and full erasure coding.
- Evaluate the recovery times for double failures.

The non-goals of this project are:

- Design a global distributed file system for Wide Area Network which is ready for a real world wide area deployment.
- Do not expect this to work for generic workload of traditional servers, or act as a replacement of any sort.
- Do not expect to be an equivalent of cloud computing environment for the Edge layer.

4. Key Contributions

The key contribution of this work is to provide a hybrid technique to store data reliably, it takes both Erasure Coding and Replication and provides best of both the techniques. For a given file, in the PRPE technique , using the given storage budget, I try to optimally store portions of file in replicated way and the rest in an erasure coded way.

This study should shed light on the trade off of compute versus storage, and show how read/write performances vary from a completely compute intesive scenario to completely data I/O scenario.

The study also examines the time to perform recovery of files using the proposed PRPE technique.

5. Assumptions

In order to achieve the goals we make the following assumptions.

- All edge devices can talk to each other.
- All edge devices can talk to the master node and vice versa.
- Master node will not fail and is hosted in one of the desktop grade computers.
- Edge devices have enough storage to host all the writes.
- All the files stored in master node are in a flat namespace.
- No appends or edits to the data.
- Edge Devices will fail.
- The edge devices all have same reliability and it is assumed to be 90
  Some of these assumptions maybe relaxed later, it will be stated explicitly if and when they are relaxed.

6. Motivation and Related Work

6.1. Motivation

IoT devices have seen a remarkable acceleration in deployment in the market in the recent years due to the practical relevance of their applications, cheaper compute costs and reduced pricing of devices.

The IoT devices which are deployed in the real-world act as sources of data (Eg: A video surveillance camera may periodically send small video streams, a temperature sensor in a laboratory might constantly stream its data). Performing fast analytics on these data are one of the key things to provide quality services in applications.

Cloud data centers are one of the solutions for reliably storing such data but, it comes with an associated storage monetery costs and data access latency.

So, as much as possible, we would like to have data close to where it is generated to gain on read latency and to avoid unwanted movement of data.

One of the major concerns of storing data at the edge is that the edge devices are low on reliability. They are often placed at outdoor environments, subjected to heat, cold and winds. The chances of devices failing are quite high.

The need for low latency and the inherent low reliability of edge device deployments provide an interesing challenge in building a reliable storage at the edge layer.

One of the most popular techniques used to store the data reliably is Replication, where reliability is achieved by storing multiple copies of data. This is a simple and elegant solution and offers reliable storage against failures, often, the data is replicated thrice to achieve the needed reliability. It is evident that the benefits come from an additional storage cost.

The alternative technique has been Erasure Coding which provide same reliability in half the storage space [1], but trades the storage savings with CPU compute. Nevertheless, it provides fault tolerance in the same level as that of Replication.

While both these techniques work quite well for commodity machines in the LAN data center, it is interesting to see how well these techniques work on the Edge devices.

The problem space is examining replication and erasure coding at the edge devices.

Replication offers the best read/write performance since they don't involve any compute during read and write, but, uses up 3x the space of the original block (considering the blocks are stored in an uncompressed manner). On the other hand, erasure coding offers same reliability in nearly 1.5x (for eg: RS(6,4) erasure coding technique) the storage but with an additional compute costs.

Since edge devices is modest on storage and compute, it is interesting to see how the proposed PRPE techniques work for storing files in Edge devices, which is what is the current study is about.

6.2. Related Work

HDFS is a distributed file system, used extensively in Facebook, Yahoo and other companies. HDFS guarantees reliability of the files it stores using replication. So as long as one copy of the file is present and accessible, the file is available. Using replication to provide reliability incurs a lot of storage cost, we essentially use up 3 times the storage space compared to a single file. To reduce this extra storage space and still provide the same reliability, it started using Erasure Coding [2].

A lot of work has been done in comparing the performance of Erasure Coding with Replication in Local Area

Network, but there's a need for a comprehensive investigation of read/write performance characteristics in edge devices.

Similarly, [3] Hyrax:Cloud Computing on Mobile Devices using MapReduce has looked into the aspects of porting the HDFS on the Android Devices, however, this hasn't looked into the aspects of erasure coding.

7. Problem Definition and Approach

Given a system comprising of a set of edge devices which are managed by a master node, *the objective is to use the collective storage of the connected Edge devices to provide reliable distributed storage for a given file with an associated storage budget using replication and erasure coding.*

Given a file, and an associated storage budget, the following equation decides how many blocks are to replicated and how many are to be erasure coded. *

$$Nr = (K/N) * SB * totalBlocks \quad (1)$$

$$Ne = totalBlocks - Nr \quad (2)$$

Where,
- $Nr \rightarrow$ number of blocks that can be replicated
- $Ne \rightarrow$ number of blocks that needs to be replicated
- $SB \rightarrow$ storage budget given by the user
- $totalblocks \rightarrow$ total number of blocks for a given file.
- $K \rightarrow$ data blocks in the given file.

Using this we can calculate the number of blocks that need to be stored using erasure coding, as shown in eqn 2.

Note: N and K are the standard parameters in the Reed Solomon encoding techniques. [4]

The erasure coding technique that I will be using throughout the study is the Reed-Solomon Erasure Coding, RS(N,K) which tolerates upto 2 nodes failure, where N is the total number of coded blocks, K is the number of data blocks

The figure 1 shows a high level architecture of the proposed system, it is a simple client server model, with one master node managing multiple edge nodes. The master node has an in-memory data structure which will help the client perform reads() and writes().

As we can see the 3 major modules are:
- Master Node
- Clients
- Edge Devices

The master node holds a mapping of files to blocks, and blocks to edge nodes locations to support replication, and some additional amount of information to perform decoding of erasure coding blocks.

The Client is the originator of reads and writes, and it interacts with master node to write and read data.

The Edge devices are storage nodes, they will have a Erasure encoding/coding module , to perform encoding and writing of incoming block and decoder module to perform decoding of blocks to return a block to the client.
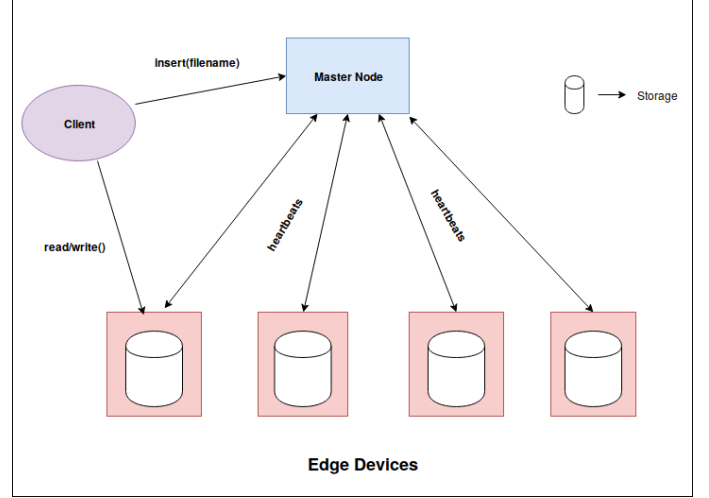

Figure 1: Master node and a set of Edge devices

7.1. Proposed Approach and The APIs to implement

The initial bootstrapping happens as follows:
- The master node is deployed on a desktop grade system on a publicy known IP.
- A configuration file is distributed to all the edge devices so that they can join the master.
- The edge devices are deployed on containers which are configured to Edge device specifications, they connect to the central master node, and send frequent hearbeats and block reports.

After this the clients can issue read and write requests to the master server.

The APIs used to solve the problem is as follows:

7.2. Apis to implement

- **Master.open(filename)** : The *open()* api is called by the client on the master node, which wants to store a file in the system. It returns a sessionid, this is used as a key for subsequent communication of client with the master node.
- **Master.put(choice)** : The *put()* api is called by the client on the master node, which specifies the choice. the choice can be either replication or erasure coding. This api returns a list of block numbers and corresponding device locations for client to write to.
- **Master.write(data,blocknumber)** : This *write()* api is called by the client on the edge devices. This is to write the data on to the edge device directly by the client.
- **Master.getBlockNumbers(filename)** : The *getBlockNumbers()* api is called by the client on the master, this will return a set of block locations.
- **Master.get(blocknumber)** : The *get()* api is called by the client on the master,This is to read the data from the edge device directly by the client.

- **Master.close(sessionid)** : The *close()* api is called by the client on the master, This api is called after writing all blocks of data. The sessionid is used to identify the client and termination session beyond time outs. This api returns SUCCESS or FAILURE. If it is SUCCESS, the metadata is persisted at the master node.

7.3. Feasibility

The CPU speed , main memory and storage of the edge devices are sufficient to run the encoder/decoder module on the edge devices. In terms of specifications of the devices it is feasible to run the data node server on it.

7.4. Hardware and Software Requirements

To conduct these evaluations we need a set of 10 containers each to run the edge devices server, and a single node to run the master node.

The hardware requirements for the master node is. (this may change later, in such circumstances it'll be duly reported)

System Environment of Master Node(Indicative):
- **CPU :** Intel(R) Core(TM) i5-8250U CPU @ 1.60GHz
- **Total Installed RAM:** 8GB
- **CPU Architecture :** 64-bit
- **Number of Processing Cores :** 8
- **Java Version:** 1.8.0201
- **Message Exchange format:** Thrift messages
- **OS :** Ubuntu 16.04 LTS

System Environment of Edge Devices (Indicative):
- **CPU :** Single Core of Intel i5-8250U Processor running at 1.6 GHz
- **Installed RAM:** 1 GB
- **CPU Architecture :** 64-bit
- **Number of Processing Cores :** 1
- **Java Version:** 1.8.0201
- **Message Exchange format:** Thrift messages
- **OS :** Ubuntu 16.04 LTS
- **Storage :** 10GB

8. Experimental Design

10 instances of edges are spawned in separate containers and master also is running in the same overlay network in a container.

8.1. Dataset for experiment

Used a file of size 10MB (10.8MB to be precise). The put and get experiment was conducted with different block sizes viz. 0.25MB, 0.5MB, 1MB, 2MB and 4MB. The put and get showed best performances (lowest delay) for 1MB. Hence the block size was chosen as 1MB.

The block size chosen is 1MB.

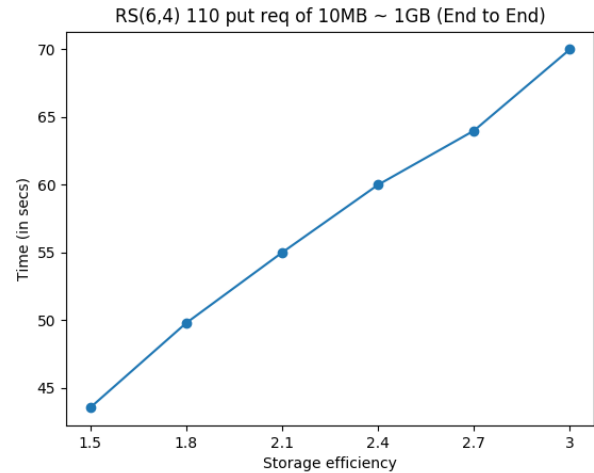The experiments were conducted as stated below:

Figure 2: 1.1GB put file

- Measure the read/write latency of putting a file in the data store with full replication/full erasure coding.
- Measure the read/write latency of putting a file in the data store with repliation-erasure coding method conduct the experiments for the following set of values 1.5x, 1.8x, 2.1x, 2.4x, 2.7x, 3.0x
- Plot the graph of how the read/write latency goes from 1.5x to 3x.
- Kill 2 nodes in the cluster and recovery every block which belonged to the killed nodes (for 1.5x to 3x).

Note: I have used Reed Soloman RS(6,4) Erasure coding technique to encode and decode the blocks in all the experiments.

8.2. Put Experiment

In this I put 100 different files of size 10MB (10.8MB) to the datastore, which is a total of 1.1GB. I varied the storage budget from 1.5X to 3.0x. For 1.5X all the files were fully erasure coded. At 3.0X all the files were fully replicated.

The time taken to write the blocks are shown in the plot 2. X-axis denotes the storage efficiency, Y-axis denotes the time in seconds.

- **Expectation**: The put for the highest storage efficiency is expected to perform best since, the amount of data being written is relatively less. Replication was expected to take the highest time.
- **Observation**: The results were as expected, since replication had to write the blocks 3 times, it took more time compared to fully erasure coded technique which took 1.5 times.

8.3. Get Experiment

In this, for all the 100 files that were put,I did a get on them. Again files which were stored with different storage

Figure 3: 1.1GB get file

Figure 4: E-Blocks vs R-Blocks

budget were read. A total of 1.1GB was read for each storage technique from 1.5x to 3x.

- **Expectation**: A get request of a file which is fully erasure coded was expected to be slow, while the file which was stored with fully replicated was expected to be fast.
- **Observation**: A fully replicated technique worked the best, since it involved only reading a block also it involved only a single read. Where as in fully replicated technique, for RS(6,4), to reconstruct a minimum of 4 blocks had to be pulled to an edge, decoded and then sent back. Hence it was slow.

The end to end time taken to read the blocks are shown in the plot 3. X-axis denotes the storage efficiency, Y-axis denotes the time in seconds.

We should read the results of put and get experiments with the change in the number of replicated blocks (R-Blocks) and Erasure Coded blocks (E-blocks). It can be seen that, when we have more and more replicated blocks, the put and get experiments will perform well.

The number of R-Blocks and E-Blocks as we vary the storage budget is shown in the plot 4

8.4. Recovery Experiment

As a part of recovery experiments, in the setup of 10 nodes, 2 nodes were randomly killed. All the blocks in the nodes that were killed are recovered completely. The recovery experiment was conducted for storage efficiency ranging from (1.5x to 3x).

- **Expectation** : Fully Replicated system was expected to recover fast, Fully Erasure coded system was expected to be the slowest.
- **Observation** : As expected, fully replicated system recovered fast since it involved just reading and writing of a block, in replicated system it involved downloading


Figure 5: Recovery of blocks from 2 dead nodes

a minimum of 4 blocks, reconstructing the file and encoding and writing the missing block.

The plot of recovery times are shown the graph below 5. X-axis denotes the storage efficiency, Y-axis denotes the time in seconds.

In the dead nodes, the split up of the number of blocks which were erasure coded and number of blocks that were replicated are shown the the plot below 6:

8.5. Erasure Coding Families

In this experiment, I wanted to see the effect of using Erasure Coding techniques which offer more efficient storage than RS(6,4) and examine how it impacts read() and write() times.

For this experiment, I have done a put() and a get(), where 100 puts of 10MB (10.8MB) file were done, a total of

Figure 6: Recovery of blocks from 2 dead nodes

Figure 7: Erasure Coding Family RS(6,4) RS(8,6) RS(10,8)

1.1GB were written and read. Two additional Reed-Soloman Codes were used to compare against RS(6,4).

The codes that were used are RS(6,4) , RS(8,6) and RS(10,8). The storage efficiencies of these respective codes are 1.5x, 1.35x and 1.25x respectively. And all are comparable in terms of fault tolerance since all the 3 techniques used support upto 2 failures.

The put() and the get() time observed are shown in the plot below 7

The put() times are comparable, all the 3 techniques took nearly 50s to store a 1.1GB in the data store, but the get() times started performing very bad since, *for each block* read a minimum of 6 and 8 blocks were read, which resulted in an addition of 2 and 4 blocks comapred to RS(6,4).

9. Challenges faced

- Handling the metadata for both replication and erasure coding per file was a tricky management.
- The recovery mechanism was very interesting in cases of a mix of coded and replicated blocks.
- This study would've been more suitable to be ran on a cluster which doesn't lack resources to spin a good number of containers.

10. Conclusion

This study has proposes a novel technique called PRPE (Partial Replication Partial Erasure Coding), which splits the file into components which can be erasure coded and which can be replicated based on a storage budget.

It investigates how the read and write latencies change with respect to change in the storage efficiency from (1.5x to 3x) and experimentally shows that erasure coding performs well on writes and relatively bad compared to replication when it comes to reads.

Also demonstrated recovery with a failure of upto 2 nodes, here too fully replicated blocks perform very well on recovery, this too has been shown experimentally.

Finally, a comparison was made between different Erasure Coding techniques in the Reed-Soloman Family, which shows that as we move towards tighter storage efficiency, the read performances become very bad.

11. Future Work

Paralellize the encoding and decoding components of Erasure coding, to improve the performance of read() and writes(). Improve the recovery speed for single node failures.

References

[1] K. Rashmi, N. B. Shah, D. Gu, H. Kuang, D. Borthakur, and K. Ramchandran, "A solution to the network challenges of data recovery in erasure-coded distributed storage systems: A study on the facebook warehouse cluster." in *HotStorage*, 2013.

[2] M. Xia, M. Saxena, M. Blaum, and D. Pease, "A tale of two erasure codes in hdfs." in *FAST*, 2015, pp. 213–226.

[3] E. E. Marinelli, "Hyrax: cloud computing on mobile devices using mapreduce," Carnegie-mellon univ Pittsburgh PA school of computer science, Tech. Rep., 2009.

[4] J. S. Plank, "Erasure codes for storage systems: A brief primer," *Login: The USENIX Magzine*, vol. 38, no. 6, pp. 44–50, 2013.